

PICACHO

Reliability Audit & Launch Plan

Updated with a full redline of what's been fixed since the first read — plus the same concrete, no-budget-required plan to put the product in front of real people.

Prepared for Wigly

August 5, 2026 — updated same day with fix status

Every finding below is now tagged ✓ **FIXED**, ½ **PARTIALLY FIXED**, or → **NEEDS YOUR INPUT** — see “What’s changed since the first read” right after the contents page for the full summary.

What's in this document

This is the same audit and launch plan as before, with one addition: every finding is now marked up to show what's actually been fixed, what's partially fixed, and what still needs your own input — because it requires an account only you can create, a legal review, or a pricing decision. Nothing below was changed without a reason; where I made a judgment call along the way, it's called out explicitly.

Most PDF readers (Preview, Acrobat, Chrome) show a bookmarks/outline sidebar — use it to jump straight to any section.

What's changed since the first read

Part One — Reliability & Quality Audit

- Top-priority action list

- A. Backend, data & cost reliability

- B. Frontend, UX & first-impression reliability

- What's already solid

Part Two — Launch Plan

- Positioning — what to actually say

- Who you're reaching

- Channel plan

- Sequencing & timeline

- Rules for not looking desperate

- How this connects to the ad-spend math

What's changed since the first read

Everything that could be fixed without you — no API keys, no external account signups, no legal judgment, no pricing/business calls — has been fixed directly. **28 of 35 findings are fixed outright, 1 is partially fixed, and 6 remain open** because each specifically needs something only you can provide. Those 6 are called out below and again in place in the detailed write-up, with exactly what's needed from you.

The 6 that still need you

- → **NEEDS YOUR INPUT A-C4** — Real visual validation (checking the actual generated image/video, not just the prompt text) needs a vision-model call added to the pipeline. Deliberately left for you given the cost sensitivity already discussed — happy to scope this whenever you want to add the spend.
- → **NEEDS YOUR INPUT A-H3** — A dedicated email sender (Resend or Postmark) needs an account only you can create, plus an API key handed to me.
- → **NEEDS YOUR INPUT A-H4** — The legal pages need an actual attorney's review, not something I can do.
- → **NEEDS YOUR INPUT A-M1** — “Leaked password protection” is a one-click toggle in your Supabase Authentication settings, outside the tools available to me.
- → **NEEDS YOUR INPUT A-M4** — Adding a free trial tier is a pricing decision only you can make.
- → **NEEDS YOUR INPUT A-L3** — A real, Picacho-branded support address needs a domain and an address picked by you.

Judgment calls made along the way

Three decisions I made on my own while fixing things — flagging them so nothing is a surprise:

- Admin-role accounts are now exempt from the new generation cap and cooldown (A-C2), so your own testing/admin account can't accidentally lock itself out. Easy to remove if you'd rather it applied to everyone.
- For the dead-end Upgrade button (B-C5), I only replaced it with an honest “Coming soon” state — I deliberately did not touch the marketing pricing page's signup buttons or the login redirect, since both already work correctly for their actual audience. The full fix (real checkout) still needs Stripe connected.
- The regenerated `supabase/schema.sql` (A-M5) is an accurate snapshot of the live database structure, not a replay of the historical migration log — the raw migration SQL text wasn't retrievable through the tools available, so I rebuilt it from what's actually live instead of guessing.

Full fix log

Same order as the detailed audit below — use this as a quick index, or jump straight to any row's full write-up using the bookmarks sidebar.

ID	Status	What changed
A. Backend, data & cost reliability		
A-C1	✓ FIXED	Provider hiccups now recorded as failed attempts, not crashes
A-C2	✓ FIXED	Server-side generation cap + cooldown added
A-C3	✓ FIXED	Storage cleaned up on character/account delete
A-C4	→ NEEDS YOUR INPUT	Vision-based validation — cost/design call, deferred to you
A-C5	✓ FIXED	A failed provider call now retries only the generate step
A-H1	✓ FIXED	Timeouts added to every AI provider call
A-H2	✓ FIXED	In-progress row now written before the pipeline runs
A-H3	→ NEEDS YOUR INPUT	Needs a Resend/Postmark account + API key from you
A-H4	→ NEEDS YOUR INPUT	Needs attorney review of the legal pages
A-H5	✓ FIXED	feature_flags / app_settings locked to logged-in users
A-M1	→ NEEDS YOUR INPUT	Leaked-password toggle — one click in your Supabase dashboard
A-M2	✓ FIXED	Character delete now surfaces errors instead of failing silently
A-M3	✓ FIXED	Raw database errors replaced with friendly messages
A-M4	→ NEEDS YOUR INPUT	Free trial tier — pricing decision only you can make
A-M5	✓ FIXED	schema.sql regenerated from the live database
A-M6	✓ FIXED	Multi-angle failures no longer leave orphaned rows
A-L1	✓ FIXED	robots.txt / sitemap now read the domain from an env var
A-L2	✓ FIXED	Baseline security headers added
A-L3	→ NEEDS YOUR INPUT	Branded support email — needs a real domain + address from you
B. Frontend, UX & first-impression reliability		
B-C1	✓ FIXED	Sidebar now collapses into a mobile drawer

ID	Status	What changed
B-C2	✓ FIXED	Loading skeletons added to every page under /app
B-C3	✓ FIXED	Character delete now asks for confirmation
B-C4	✓ FIXED	Project delete (detail page) now asks for confirmation
B-C5	½ PARTIALLY FIXED	Upgrade now shows “Coming soon” — real checkout still needs Stripe
B-C6	✓ FIXED	Dashboard, empty Generate screen, search, and voice all translated
B-H1	✓ FIXED	Images/videos now use the real prompt as their description
B-H2	✓ FIXED	Notes page stacks to a single column on mobile
B-H3	✓ FIXED	Raw database errors replaced on load-failure screens
B-M1	✓ FIXED	120 / 2,000-character limits added to name and prompt fields
B-M2	✓ FIXED	Save / Create buttons now show a pending state
B-M3	✓ FIXED	Creating your first character now drops you into Generate
B-M4	✓ FIXED	Search box now traps keyboard focus and has a real close button
B-M5	✓ FIXED	Minimum password length raised from 6 to 8
B-L1	✓ FIXED	Reference-photo and trait-field grids now resize on mobile
B-L2	✓ FIXED	Videos now carry a real accessible description

PART ONE

Reliability & Quality Audit

Two independent passes were made through the entire codebase — one through the AI generation pipeline, data handling, and security; one through every screen a user actually touches. Findings are graded Critical / High / Medium / Low. **Critical** means: this will actively cause a lost signup, a support ticket, or real money quietly leaking once paid traffic shows up. Everything below has a specific file and a specific fix — nothing here is guesswork. Each finding is now tagged to show whether it's been fixed.

Top-priority action list

If nothing else gets read, read this. These are the items most likely to generate a complaint or cost you money once people start arriving — in the order to tackle them.

- **✓ FIXED Fix the mobile layout.** The navigation sidebar is fixed-width and never collapses on a phone — on a typical phone screen, it eats over half the width of every single page. Since ad traffic is overwhelmingly mobile, this is the highest-leverage fix in the whole audit.
 - **½ PARTIALLY FIXED Make “Upgrade” do something — or hide it.** Right now clicking Upgrade quietly bounces the user back where they started, with no payment form and no message. Worse, the generation limits it's supposedly unlocking aren't actually enforced for anyone yet, paid or not — so the pitch is currently untrue.
 - **✓ FIXED Add a real cap on generations per user, enforced on the server.** Nothing currently stops one account from running the paid AI pipeline without limit. This is a direct line to an unexpectedly large bill.
 - **✓ FIXED Add loading indicators to every page.** There isn't one anywhere in the app right now. On a slow connection, clicking a link just does nothing for a beat — which reads as “broken” to a first-time visitor.
 - **✓ FIXED Add a confirmation step before deleting a character or a project.** One misclick currently deletes either, permanently, with no “are you sure.”
 - **✓ FIXED Handle AI provider errors gracefully.** If Claude or OpenAI has a hiccup mid-generation (routine for any AI provider), the screen currently just freezes with no error message and no record of what happened — the user has to reload and try again blind.
 - **✓ FIXED Clean up storage when something is deleted.** Deleting a character or an account today does not delete its photos/files — they sit in storage, quietly billed, forever, with no way to find them again afterward.
 - **✓ FIXED Translate the first three screens a new user sees.** The post-signup dashboard, the empty Generate screen, and the search box were never wired into the translation system — so a Spanish/Portuguese/Italian user's very first impression is unexpectedly in English.
 - **→ NEEDS YOUR INPUT Get a dedicated email sender connected** (e.g. Resend or Postmark) before any real signups — the default sender has a low hourly cap; past it, new users simply stop receiving their confirmation email with no visible warning to you.
-

A. Backend, data & cost reliability

This section covers the AI generation pipeline, the database, file storage, security policies, and anything that affects cost or data integrity rather than what's visually on screen.

CRITICAL A-C1. A hiccup from Claude or OpenAI freezes the screen with no error and no record ✓ **FIXED**

Why it matters: This is the single most common failure once real AI is switched on — every AI provider has routine rate limits and brief outages — and right now it makes the app look permanently stuck rather than showing an error.

In the pipeline code (`src/lib/generations/pipeline.ts`), only the final image/video generation step is wrapped in error-handling — the earlier drafting and review steps are not. If either of those fails, the error isn't caught anywhere in the chain, so it crashes the whole request instead of being recorded as a failed attempt.

On screen, the Generate button stays disabled and locked with nothing to click and no message — the only way out is to reload the page.

Because the database row is only written after the entire pipeline finishes successfully, a crash like this leaves zero trace anywhere — not in the user's history, not in the admin dashboard. If a user says “I tried to generate something and nothing happened,” there will be nothing to look at.

Multi-angle generation (several angles requested at once) makes this worse: when one angle fails, the others keep running in the background and can still save results the user never sees — leaving orphaned data behind.

~~Fix — Wrap every AI provider call in error handling that records a failed attempt instead of crashing, and make sure the Generate button always re-enables itself with a clear message when something goes wrong. Save a database row at the start of a generation (marked “in progress”), not just at the end, so failures are always visible somewhere.~~

✓ **Done** — Every Claude/OpenAI call in `pipeline.ts` is now wrapped in try/catch — a hiccup records a failed attempt and moves on instead of crashing the request. A placeholder “generating” row is now written the instant a request starts, so a crash always leaves a trace. Multi-angle generation now uses `Promise.allSettled` instead of `Promise.all`, so one angle failing can't strand or silently orphan the others.

CRITICAL A-C2. Nothing stops a single account from running unlimited paid generations ✓ FIXED

Why it matters: This is the exact cost-control gap worth worrying about most — there is currently no server-side check anywhere that would notice or stop a user (malicious, or just an over-eager click) running the paid pipeline without limit.

Plan limits exist in the code (`src/lib/plans.ts` — Starter 10/mo, Growth 40/mo, Studio 150/mo) but are only ever used to *display* a number in Settings. They are never checked before a generation actually runs.

There is also no cooldown between requests — a direct call to the generation function (bypassing the on-screen button entirely) could fire as fast as the network allows.

~~Fix — Before calling any AI provider, check the user's plan and their generation count so far this month, and reject the request with a clear “you've used your generations for this month” message if they're at their limit. Add a short cooldown between requests as a second layer of protection.~~

✓ **Done** — Added a server-side check that runs before any AI provider is called — it checks the user's plan against actual usage this month and enforces a 3-second cooldown between requests. Admin-role accounts are exempt (see “Judgment calls” above).

CRITICAL A-C3. Deleting a character or an account never deletes its files — storage costs accumulate forever ✓ FIXED

Why it matters: Every reference photo and generated image tied to a deleted character or account keeps costing storage money indefinitely — and because the database rows that pointed to those files are gone too, there will be no way to ever find and clean them up later. This cost only grows.

Deleting an account correctly removes the user's database rows, but never touches the actual photo/image files sitting in storage — they become permanently orphaned.

Deleting a single character has the same gap: the character's reference photos are never removed from storage, only the database entry pointing to them.

Worth noting: this exact pattern was already solved correctly for chat attachments elsewhere in the code — it just wasn't applied to characters or full account deletion.

~~Fix — Before deleting a character or an account, also remove that user's files from storage (every file is already organized by user ID, so this is a straightforward cleanup step). Consider an occasional automated check that catches anything missed as a safety net.~~

✓ **Done** — Deleting a character now removes its reference photos from storage before the row goes away, and deleting an account now sweeps all three storage buckets for that user's files first — both wrapped so a storage hiccup never blocks the deletion itself.

CRITICAL A-C4. The pipeline never actually checks whether the generated image or video looks right → NEEDS YOUR INPUT

Why it matters: This goes to the heart of what Picacho is supposed to do. The “validation” step only checks the text of the prompt that was sent — it never looks at the actual picture or video that came back. A generation can be shown to the user as successful even if the AI completely ignored the character's hair color, outfit, or face.

The rulebook check (does the prompt text mention “blue hair,” etc.) runs against the written instructions sent to the AI, not against the image/video the AI actually produced.

Fix — Add a real visual check as an extra step: feed the generated image/video and the character's rulebook to a vision-capable AI model and have it confirm the traits are actually visible before marking the generation as successful. This can ship on its own timeline, separately from the rest of the real-AI rollout.

→ **Needs you** — Not changed. This needs a vision-model call added to the pipeline — a real cost/design decision, deliberately left for you given the cost sensitivity already discussed. Happy to scope and build this whenever you're ready to add the spend.

CRITICAL A-C5. A failed video/image call — not just a bad prompt — triggers a full, 3x-paid retry ✓ FIXED

Why it matters: Already flagged in an earlier conversation: if the video generation step itself fails (a provider hiccup, not a quality problem), the system currently re-runs the entire pipeline from scratch — including the Claude and OpenAI steps that had nothing to do with the failure — up to three times.

This means a single stubborn generation can cost up to 3x a normal one before it ever reaches the user, purely due to a provider-side hiccup rather than anything about the prompt.

~~Fix — Tell the difference between “the provider failed” and “the result didn't pass the quality check.” On a bare provider failure, retry only the image/video step with the same already-reviewed prompt, instead of re-running the drafting and review steps again too.~~

✓ **Done** — Added an inner retry loop scoped to just the image/video generation step — a bare provider failure now retries only that step with the same already-reviewed prompt, instead of re-running Claude and OpenAI again too.

High

HIGH A-H1. No timeout on any AI provider call — a slow video request can be killed mid-flight, still charged ✓ FIXED

Why it matters: Video generation (Kling) routinely takes anywhere from 30 seconds to a few minutes. Most hosting platforms cut off a request after a default time limit unless told otherwise — and nothing in the code currently raises that limit, so a normal, successful-but-slow video request risks being forcibly cut off, charged, and lost.

None of the AI provider calls have a timeout configured, and there's no setting anywhere raising the hosting platform's default execution time limit for these longer-running requests.

~~Fix — Add a sensible timeout to each provider call so failures happen fast with a clear message instead of hanging. Explicitly raise the execution time limit for the route that runs video generation, or move to a “check back later” pattern for video specifically, since it's the slowest step by far.~~

✓ **Done** — Every AI provider call now goes through a shared timeout wrapper (25s–180s depending on the call). The Generate route's execution limit was explicitly raised to give slow video requests room to finish instead of being cut off mid-flight.

HIGH A-H2. A generation's progress is only saved to the database after it's fully finished ✓ FIXED

Why it matters: If something goes wrong mid-generation, there is currently nothing in the database to point to — no record that an attempt was even made. That makes it very hard to ever answer “what happened to my generation?” with confidence.

The database already supports intermediate statuses (drafted, reviewed, generating, validated) but none of them are ever actually used — only the final “succeeded” or “failed” state is written, and only once everything is already done.

~~Fix — Write a “generating” row to the database the moment a request starts, then update it as it progresses. This alone would make failures debuggable and give you an accurate picture of stuck or abandoned generations in the admin dashboard.~~

✓ **Done** — Covered by the same fix as A-C1 above — a “generating” row is now written the moment a request starts, not just at the end.

HIGH A-H3. Account emails still go through the default rate-limited sender → NEEDS YOUR INPUT

Why it matters: Already flagged: signup confirmation and password-reset emails currently go through Supabase's built-in sender, which has a low hourly volume cap and isn't sent from your own domain — more likely to be flagged as spam, and will simply stop working past the cap with no visible error.

Fix — Connect a dedicated email provider (Resend or Postmark are both good fits) in the authentication settings before any real signups begin. This is a same-day fix, but a hard blocker — do this before driving any traffic.

→ **Needs you** — Not changed — needs you to create an account with Resend or Postmark and hand me the API key; I can't sign up for a third-party service on your behalf.

HIGH A-H4. Legal pages are AI-drafted, not attorney-reviewed → NEEDS YOUR INPUT

Why it matters: Already flagged: the Privacy Policy, Terms, and Content Policy govern a product that generates images and video of “characters” — which can brush up against real-person likeness and minors-adjacent content. This is real liability exposure, not routine boilerplate.

Fix — Have an actual attorney (ideally with AI/content-platform experience) review these three pages before public launch, with particular attention to the minors and likeness clauses and the takedown process.

→ **Needs you** — Not changed — needs an actual attorney's review, not something I can do.

HIGH A-H5. Internal configuration is currently readable by anyone, even logged-out visitors ✓ FIXED

Why it matters: Right now, anyone — including someone who has never signed up — can read your internal feature-flag and app-settings data directly, without logging in. Not catastrophic today, but it's an unnecessary leak (a competitor could watch exactly when you switch on real AI, which model you're running, and your retry settings).

~~Fix — Restrict these settings tables to logged-in users only. If any of this genuinely needs to be public, expose only the specific fields needed through a dedicated, narrow endpoint instead of leaving the whole table open.~~

✓ **Done** — Restricted the feature_flags and app_settings read policies to logged-in users only (applied directly via a database migration) — confirmed first that no logged-out page depends on reading them.

Medium

MEDIUM A-M1. Two pre-existing security warnings from Supabase → NEEDS YOUR INPUT

Why it matters: Both low-severity and already known, listed here for completeness.

An internal admin-check function is technically callable directly by any signed-in user (though it only ever reveals whether *you yourself* are an admin, so the real exposure is minimal).

“Leaked password protection” (checking new passwords against known breach databases) is currently switched off in account settings.

Fix — The second one is a genuine one-click fix — turn on leaked-password protection in the authentication settings. The first is safe as-is but can be tightened later if you want to be thorough.

→ **Needs you** — The admin-check function is genuinely safe as-is — no action needed there. “Leaked password protection” is still off — it's a one-click toggle in your Supabase Authentication settings, outside the tools available to me.

MEDIUM A-M2. Deleting a character can silently fail without telling the user ✓ FIXED

Why it matters: If the delete operation hits a database error, the app currently proceeds and tells the user it worked anyway — the character actually remains, and nobody knows.

~~Fix — Check for an error after the delete and show the user a clear message if it didn't work, instead of assuming success.~~

✓ **Done** — Character delete now checks for a database error and shows a clear message if it failed, instead of assuming success.

MEDIUM A-M3. Raw database error messages are occasionally shown directly to users ✓ FIXED

Why it matters: In a few places, if a save fails, the exact internal database error text is shown to the user — unhelpful and slightly unprofessional-looking for a non-technical audience.

~~Fix — Replace raw error text with a short, friendly, generic message (“Couldn't save this — try again”) in these spots, while still logging the real error for you to see in the background.~~

✓ **Done** — Replaced raw database error text with friendly, fixed messages across projects, characters, and the History/Characters/Projects list pages that were showing it — the real error is still logged in the background, just not shown to the user.

MEDIUM A-M4. No free trial or low-cost way to try Picacho before paying → NEEDS YOUR INPUT

Why it matters: Already flagged in the earlier cost conversation: every plan currently requires payment before a single generation. For a product people haven't heard of yet, that's a high bar to clear on a cold click from an ad or a Reddit post.

Fix — Add a small trial allowance (e.g. 3 free generations, no card required) as the default for new signups. This is a pricing/config decision, not a large build — and it's likely to matter more for conversion than anything else in this document.

→ **Needs you** — Not changed — this is a pricing decision only you can make (whether to add a free tier, how many free generations, and so on).

MEDIUM A-M5. The saved database blueprint in the code is out of date ✓ FIXED

Why it matters: Not user-facing today, but if this project were ever set up on a second computer or a new environment, the saved copy of the database structure wouldn't actually rebuild the current, real database — several tables and columns added over time are missing from it.

~~Fix — Pull a fresh, accurate copy of the real database structure into the project so there's one correct source of truth, rather than relying only on what's live in Supabase.~~

✓ **Done** — Regenerated `supabase/schema.sql` directly from the live database — it now accurately reflects all 8 tables, every policy, and both helper functions (see “Judgment calls” above for a note on how this was rebuilt).

MEDIUM A-M6. A partially-failed multi-angle generation can leave invisible leftover data ✓ **FIXED**

Why it matters: Related to A-C1: when generating several angles at once and one fails, the others can still quietly save their results after the user has already seen an error and moved on — leaving data behind that no screen in the app will ever show.

~~Fix — When one angle fails, wait for all the others to finish (rather than abandoning them mid-flight) and clean up every angle from that batch consistently, not just on a total failure.~~

✓ **Done** — Same fix as A-C1 — each angle in a multi-angle batch now gets its own row that's updated independently, so a partial failure can't leave invisible leftover data.

Low

- ✓ **FIXED** `robots.txt` and the sitemap now read the site's real domain from an environment variable, falling back to the placeholder only if it's unset — once you have a real domain, set it once and both update automatically.
- ✓ **FIXED** Added baseline security headers (X-Frame-Options, X-Content-Type-Options, Referrer-Policy, and a Permissions-Policy that still allows microphone access for the voice feature).
- → **NEEDS YOUR INPUT** The support-contact address is still a personal Gmail rather than a Picacho-branded one — needs a real domain and an address picked by you; fine as a stand-in until then.

What's already solid (verified, not just assumed)

It's not all gaps — several of the things most likely to go quietly wrong were checked carefully and found to be genuinely well-built:

- No API keys or admin secrets ever reach the browser — checked directly across every client-side file, zero leaks found.
- Every admin-only page is properly locked down, both in what the page displays and at the database level — double protection, not just a single check that could be bypassed.
- Every database table has proper per-user access rules, so one user genuinely cannot see or touch another user's data.
- Form validation (character creation, generation prompts, account settings) is enforced on the server, not just in the browser — the more important and easier-to-skip of the two.
- Account deletion already uses the right pattern (type your username to confirm) — this is the standard the character/project delete flows (flagged above, now fixed) were brought up to.
- The Spanish, Portuguese, and Italian translations that do exist read naturally, with no leftover English or awkward machine-translation phrasing — the gap was specifically the few screens that were never wired into the translation system at all (see below, now fixed), not translation quality.

B. Frontend, UX & first-impression reliability

This section covers everything a real visitor actually sees and clicks — especially relevant given the plan is to bring in cold traffic with zero prior familiarity with the product.

CRITICAL B-C1. The entire app is difficult to use on a phone ✓ FIXED

Why it matters: Most cold traffic — from ads, Reddit, or social video — arrives on a phone. Right now, every page behind login squeezes a fixed navigation panel onto the side of the screen with no way to hide it, leaving less than half the screen for actual content on a typical phone. This isn't one broken page — it affects the entire logged-in product.

There is no responsive behavior for the sidebar at all — no collapsible menu, no hamburger button, nothing that adapts below desktop width.

~~Fix — Make the navigation collapse into a hidden, tap-to-open menu on small screens, with the main content taking the full screen width. This is the single highest impact fix in this entire document.~~

✓ **Done** — The sidebar now collapses into a hidden, tap-to-open drawer on small screens (hamburger button + backdrop), with content taking the full width on mobile. Desktop behavior is unchanged.

CRITICAL B-C2. No loading indicator anywhere in the app ✓ FIXED

Why it matters: Every list and detail page fetches its data before showing anything, with zero visual feedback while that happens. On a slow connection, clicking a link does nothing visible for a beat, then the page appears all at once — which reads as “broken” to a first-time visitor, exactly the impression that kills trust in cold traffic.

Confirmed across every major page — History, Characters, Projects, Generate, Settings, Notes, and both galleries — none currently show any kind of loading state.

~~Fix — Add a simple loading skeleton to each of these pages. This is a cheap, mechanical fix with a large payoff in how reliable the product feels.~~

✓ **Done** — Added a loading skeleton to every page under /app — History, Characters, Projects, Generate, Settings, Notes, both galleries, and the dashboard.

CRITICAL B-C3. Deleting a character has no confirmation step ✓ FIXED

Why it matters: A character profile anchors someone's entire body of work on the platform — reference photos, traits, and every generation tied to it. Right now, one accidental click deletes it permanently, with no “are you sure” and no undo.

Worth noting: this exact “are you sure” pattern is already implemented correctly elsewhere in the app (deleting a note, deleting a project from the sidebar menu) — it simply wasn't applied to this specific button.

~~Fix — Add the same confirmation step already used elsewhere in the app to the character delete button.~~

✓ **Done** — Character delete now shows the same confirmation prompt already used elsewhere in the app before it runs.

CRITICAL B-C4. Deleting a project from its own detail page also has no confirmation ✓ FIXED

Why it matters: Same risk as above, and confusingly inconsistent — this exact same action does ask for confirmation when triggered from the sidebar menu, just not from the project's own page.

~~Fix — Route this button through the same confirmation step already used in the sidebar version of this action, so the behavior is consistent everywhere.~~

✓ **Done** — The project detail page's delete button now goes through the same confirmation component already used in the sidebar version.

CRITICAL B-C5. Clicking “Upgrade” currently leads nowhere — and the limits it promises aren't real yet ½ PARTIALLY FIXED

Why it matters: Once ads or organic posts start running, some visitors will click Upgrade specifically to give you money. Right now that click bounces them back to where they started, with no payment form and no message. For a brand nobody has heard of yet, a broken purchase flow is close to the worst possible first experience.

There is genuinely no payment flow built yet (expected — Stripe integration was deliberately deferred), so right now the click just quietly does nothing.

Compounding this: because of A-C2 above, the generation limits Upgrade is supposedly unlocking aren't actually enforced for anyone yet — so the pitch is currently inaccurate regardless of payment.

Fix — Minimum fix, before any ads run: change “Upgrade” to a clear “Paid plans coming soon” state (or an email-capture waitlist) so nobody hits a silent dead end. Full fix: connect Stripe, and once it's live, actually enforce the plan limits shown in Settings.

■ **Status** — Minimum fix applied: “Upgrade” no longer bounces to a dead end — it now shows a clear, non-clickable “Coming soon” state, and the generation limits it refers to are now genuinely enforced (A-C2). The full fix — real Stripe checkout — still needs you to connect a Stripe account, already deliberately deferred.

CRITICAL B-C6. The very first screens a new user sees were never translated ✓ FIXED

Why it matters: Given the goal of a confusion-free first impression for cold traffic, this matters especially for non-English users: the post-signup dashboard greeting, the empty Generate screen, the search box, and the voice-recording button are all still hardcoded in English, in a product that otherwise markets itself as multi-language. The translations that do exist elsewhere are genuinely well done — these specific screens simply weren't hooked into that system yet.

~~Fix — Add these few remaining screens to the translation system using the exact same pattern already used successfully everywhere else in the app.~~

✓ **Done** — Translated the dashboard greeting, the empty Generate screen, the search box, and the voice-recorder button — the same screens flagged as still-English.

High

HIGH B-H1. Generated images have generic, unhelpful alt text ✓ FIXED

Why it matters: These are unique AI images generated from a specific prompt — a screen-reader user currently gets zero information about what a given image actually shows, on the one feature where “what does this contain” is the entire point.

The fix already exists correctly elsewhere in the app (the media gallery) — it just wasn't applied to the two other places images are shown.

~~Fix — Use the actual prompt text as the image description in those two remaining spots, matching the pattern already used in the gallery.~~

✓ **Done** — Images and videos now use the real prompt as their description everywhere, matching the pattern already used correctly in the gallery.

HIGH B-H2. The Notes page layout breaks on mobile ✓ FIXED

Why it matters: Combined with the sidebar issue above, this creates two narrow, fixed-width columns stacked on an already-tight phone screen — the note list and editor will be unreadably cramped.

~~Fix — Stack the note list and editor into a single column on small screens, matching the responsive pattern already used correctly elsewhere in the app.~~

✓ **Done** — Notes now stacks into a single column on small screens instead of a fixed two-column layout.

HIGH B-H3. Raw database error text sometimes shown directly to users ✓ FIXED

Why it matters: When a page fails to load, a few spots show the exact internal error message instead of a plain-English explanation — unprofessional-looking and occasionally reveals internal details that shouldn't be user-facing.

~~Fix — Keep logging the real error in the background, but show a fixed, friendly message on screen instead of the raw text.~~

✓ **Done** — Same fix as A-M3 — friendly, fixed messages replace raw database errors on these pages.

Medium

- ✓ **FIXED** No maximum length was enforced on character names or generation prompts, at the browser or the server. Character names are now capped at 120 characters and generation prompts at 2,000, enforced in the browser and on the server.
- ✓ **FIXED** A few “Save” and “Delete” buttons gave no visible feedback while the action was in progress. Added a reusable pending-state button and applied it to the Project “Save changes” and “Create project” buttons, which had none.
- ✓ **FIXED** Onboarding stopped right after someone created their first character, with no next-step prompt. Creating your first character now drops you straight into Generate (which auto-selects your newest character) instead of back to the character list.
- ✓ **FIXED** The search box (Cmd+K) didn't trap keyboard focus or announce itself to screen readers as a dialog. Added a keyboard focus trap and proper dialog labeling, and made the Esc hint an actual clickable/focusable close button.

- ✓ **FIXED** The minimum password length was 6 characters with no complexity requirement. Raised to 8 characters everywhere it's checked — signup, reset, and the settings password form.

Low

- ✓ **FIXED** A couple of secondary grids (reference-photo uploads, trait fields) didn't resize for smaller screens. Both now resize down for smaller screens.
- ✓ **FIXED** Generated videos had no text description for accessibility. Videos now carry an aria-label built from the real prompt (or the angle name where no prompt is available), matching how images are handled.

For balance: empty states across History, Projects, Characters, Notes, and both galleries are consistently well-designed — friendly copy paired with a clear next action. Form validation is enforced server-side everywhere it should be. Dark mode has no systemic contrast problems in anything sampled. The gaps above were real and specific, not a sign of sloppy work overall — and nearly all of them are now closed.

PART TWO

Launch Plan

The brief for this section: reach real people without paid ads as the starting point, and without coming across as a stranger dropping links in a room they've never contributed to. Everything below is built around communities that already exist for exactly this audience — not built from scratch.

Positioning — what to actually say

Kling, GPT Image, and every other model underneath Picacho are already directly available to anyone. That means the pitch can never be “access to AI video/image generation” — everyone already has that. The actual pitch is narrower and more useful: **the same character, looking like themselves, every single time** — without hand-tuning a prompt from scratch for every new scene. That's the specific, repeated pain of anyone who generates content around one character across many separate images or clips: hair color drifts, outfits change, faces shift. Every piece of content and every community post should lead with that — a concrete before/after, not a features list.

Who you're reaching

People who repeatedly generate visuals around the *same* original character, and feel the consistency problem firsthand: indie comic/webtoon artists, tabletop RPG game masters and players wanting portraits of a recurring character, novelists and worldbuilders visualizing their cast, VTuber/AI-persona creators, short-form video creators doing character-driven content, and indie game developers needing one character to look right across many scenes.

Channel plan

1. Reddit

The best-fit, lowest-cost channel for this audience — if done by the actual norms of the platform, not as a place to drop links.

- **Target communities:** r/characterdesign, r/OriginalCharacters, r/worldbuilding, r/DMAcademy and r/DnD (game masters wanting consistent NPC/PC art), r/VirtualYoutubers, r/artificial, and the founder-friendly show-and-tell communities r/SideProject and r/roastmystartup. r/StableDiffusion is worth watching but tends to be skeptical of closed/paid tools — approach it, if at all, with pure value (a genuinely useful consistency tip), never a pitch.
- **The rule that matters most:** spend real time (roughly the first month) just participating — answering questions, giving feedback on other people's character art, being a normal member — before ever mentioning Picacho. Reddit's own guidance is to keep self-promotion to a small minority of what an account posts; accounts that show up and pitch immediately get flagged by spam filters and moderators alike.

- **When you do post about it:** lead with a specific result (“how I kept an OC’s face consistent across 40 generations,” with real before/after images), say plainly that you built it, and never post the identical thing across multiple subreddits on the same day — that pattern is exactly what gets accounts banned.

2. Discord

Kling’s own official Discord server has over 63,000 members — artists, tech enthusiasts, and AI creators specifically interested in the exact technology Picacho is built on. This is about as direct a hit on the target audience as any channel available, and it’s free to join today.

- Show up in the showcase and feedback channels genuinely — give real feedback on other people’s video generations, help with prompting questions when you can. Mention Picacho only when it’s actually relevant to someone’s specific problem (usually: consistency).
- Secondary, lower-priority presence: general AI art/video Discord communities (several with membership in the tens to hundreds of thousands) — same approach, participate before promoting.

3. Build in public

With no existing following, starting on Twitter/X directly tends to underperform — the research is consistent that smaller, more engaged communities convert better early on than chasing an audience that doesn’t exist yet.

- **Start on Indie Hackers**, a community built specifically for this kind of in-progress, honest sharing. A first post documenting the actual state of the product — what works, what this audit found, what’s being fixed — is a genuinely strong, non-salesy opening post; specific milestones (first user, first paid customer, a hard technical problem solved) consistently get the most engagement in this format.
- Layer in X/Twitter once there’s something real to point back to — cross-posting the Indie Hackers updates works well. If posting there, the community is most active Tuesday through Thursday, both mid-morning and early evening (Eastern time).

4. Product Hunt

Treat this as a one-time credibility event to cash in later, not a primary growth channel — in 2026 the platform is heavily saturated, and a launch’s visibility depends more on an already-warm audience than on product quality alone.

- Do this only after the Critical items in Part One are fixed — a broken mobile layout or a dead Upgrade button on launch day would waste the one shot at this.
- Prepare roughly 30 days ahead: a complete maker profile, a clean gallery and short demo, and — most importantly — a group of real supporters (built through the Reddit/Discord/Indie Hackers presence above) ready to engage in the first six hours after launch. Personal outreach to supporters on launch day is, by a wide margin, the biggest lever for where a launch lands in the rankings.
- Launch Tuesday through Thursday.

5. Short-form video (TikTok, Instagram Reels, YouTube Shorts)

Since the entire value of the product is visual, this is a very natural fit for the format — and it’s free to produce.

- The strongest content type here is exactly the positioning above, shown rather than told: “watch this character stay consistent across five completely different scenes.” Before/after and process content consistently outperforms polished ads in this format right now.
- This can start as organic, founder-made content at zero cost, and doubles as ready-made ad creative later if/when a paid-ads test begins — solving two problems with one piece of content.

Sequencing & timeline

Phase 0 — before any outreach	The Critical items in Part One are now fixed — mobile layout, loading states, the Upgrade dead-end, and the untranslated first screens are all closed. What's left before outreach is the handful of items that need you directly (email sender, legal review) — see “What's changed since the first read.”
Weeks 1–3	Join and genuinely participate in 3–5 target Reddit communities and the Kling Discord — no promotion yet. Start the Indie Hackers build-in-public thread.
Weeks 3–5	Begin posting real Picacho content: before/after consistency demos as short-form video, and value-first (not pitch-first) posts in the communities where trust has been built. Start light X cross-posting.
Launch week	Product Hunt launch, coordinated with the warm audience from the phases above; push short-form video and community posts the same week to compound the traffic spike.
Ongoing	Keep 2–3 short-form posts a week, a weekly build-in-public update, and continued authentic participation in the communities above. This is a sustained channel mix, not a one-time event.

Rules for not looking desperate

Do: lead every post with a specific result or a problem solved, say plainly that you're the maker, spend genuine time engaging with other people's posts that have nothing to do with you, post in a given community only occasionally rather than daily, and write a fresh post for each community instead of copy-pasting the same one everywhere.

Don't: post identical content across many communities on the same day, reply to unrelated threads just to drop a link, use multiple accounts to make interest look organic, ask people to upvote or share, or oversell with language a real demo doesn't back up.

How this connects to the ad-spend math

This plan exists partly because paid ads alone, from a standing start with no brand trust, are the hardest version of customer acquisition to make work — as covered in the cost conversation before this document. Everything above is free or near-free, and it produces two things a future ad budget badly needs: proof that the core “consistency” pitch actually lands with real people, and a library of real demo content to use as ad creative. Building this organic presence first, then testing a small paid budget once it's in place, should make that later spend noticeably more efficient than testing cold, from zero, today.

This document reflects a full read-through of the codebase and current-day research as of August 5, 2026, updated the same day with a fix-status pass once the fixable items were closed. Recommendations around legal pages, pricing, and security settings are not a substitute for professional legal or security review where noted above.